

The background of the slide is a complex, abstract network diagram. It consists of numerous nodes of varying sizes and colors (black, white, blue, and grey) connected by thin, light grey lines. Some nodes are highlighted with larger, concentric circles. The overall aesthetic is technical and modern.

Colorizing black and white images

Philip Bramwell

Agenda

- Introduction
- Project Goals
- Data
- Methods and Models
- Results

Introduction

- Using a neural network to predict RGB values
- Grayscale images as input and RGB images as labels



Project Goals

- Predict acceptable RGB images
- Use the network to predict RGB images for old black and white images that have a similarity to the training/validation data used

Data

- Mainly images of landscapes
- Relatively small dataset, so image augmentation was used
- Although the images mainly show landscapes, there are a few outliers

Data (landscape examples)



Data (outlier example)



Used Methods/Models

- ImageDataGenerator was used to provide the network with training/validation input and labels
- For the training input and label images image augmentation was used

Used Methods/Models

- Baseline model

```
input_layer = Input(shape=input_shape)

# Encoding

conv_1 = Conv2D(64, (3, 3), padding='same', activation='relu')(input_layer)
max_pooling_1 = MaxPooling2D((2, 2))(conv_1)

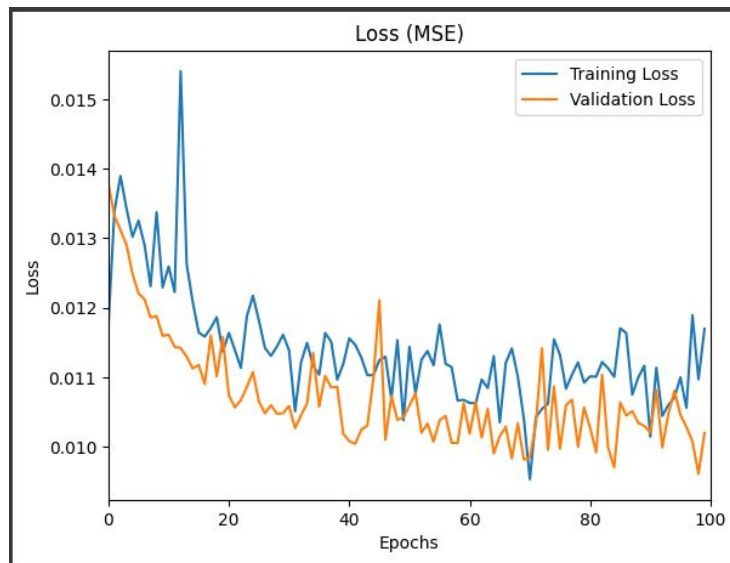
# Decoding

conv_out = Conv2DTranspose(3, (3, 3), strides=(2, 2), padding='same')(max_pooling_1)

model = tf.keras.Model(inputs=input_layer, outputs=conv_out)
```



Used Methods/Models

- The baseline model is already able to predict RGB images but the course of the loss is very unstable



Used Methods/Models

- The predicted RGB images of the baseline model are very blurry (no skip connections)



Used Methods/Models

- U-Net model

Encoding

```
conv_1 = Conv2D(64, (3, 3), padding='same')(input_layer)
batch_norm_1 = BatchNormalization()(conv_1)
relu_1 = ReLU()(batch_norm_1)
max_pooling_1 = MaxPooling2D((2, 2))(relu_1)

conv_2 = Conv2D(128, (3, 3), padding='same')(max_pooling_1)
batch_norm_2 = BatchNormalization()(conv_2)
relu_2 = ReLU()(batch_norm_2)
max_pooling_2 = MaxPooling2D((2, 2))(relu_2)

conv_3 = Conv2D(256, (3, 3), padding='same')(max_pooling_2)
batch_norm_3 = BatchNormalization()(conv_3)
relu_3 = ReLU()(batch_norm_3)
max_pooling_3 = MaxPooling2D((2, 2))(relu_3)

conv_4 = Conv2D(512, (3, 3), padding='same')(max_pooling_3)
batch_norm_4 = BatchNormalization()(conv_4)
relu_4 = ReLU()(batch_norm_4)
max_pooling_4 = MaxPooling2D((2, 2))(relu_4)

conv_5 = Conv2D(1024, (3, 3), padding='same', activation='relu')(max_pooling_4)
batch_norm_5 = BatchNormalization()(conv_5)
relu_5 = ReLU()(batch_norm_5)
```

Decoding

```
conv_trans_5 = Conv2DTranspose(512, (3, 3), strides=(2, 2), padding='same')(relu_5)
batch_norm_5 = BatchNormalization()(conv_trans_5)
relu_up_5 = ReLU()(batch_norm_5)

concat_4 = Concatenate()([relu_up_5, relu_4])
conv_trans_4 = Conv2DTranspose(256, (3, 3), strides=(2, 2), padding='same')(concat_4)
batch_norm_up_4 = BatchNormalization()(conv_trans_4)
relu_up_4 = ReLU()(batch_norm_up_4)

concat_3 = Concatenate()([relu_up_4, relu_3])
conv_trans_3 = Conv2DTranspose(128, (3, 3), strides=(2, 2), padding='same')(concat_3)
batch_norm_up_3 = BatchNormalization()(conv_trans_3)
relu_up_3 = ReLU()(batch_norm_up_3)


concat_2 = Concatenate()([relu_up_3, relu_2])
conv_trans_2 = Conv2DTranspose(64, (3, 3), strides=(2, 2), padding='same')(concat_2)
batch_norm_up_2 = BatchNormalization()(conv_trans_2)
relu_up_2 = ReLU()(batch_norm_up_2)

concat_1 = Concatenate()([relu_up_2, relu_1])
conv_out = Conv2D(3, (1, 1), padding='same', activation='sigmoid')(concat_1)

model = tf.keras.Model(inputs=input_layer, outputs=conv_out)
```


Used Methods/Models

- Learning rate 0.01




Used Methods/Models

- Learning rate 0.001



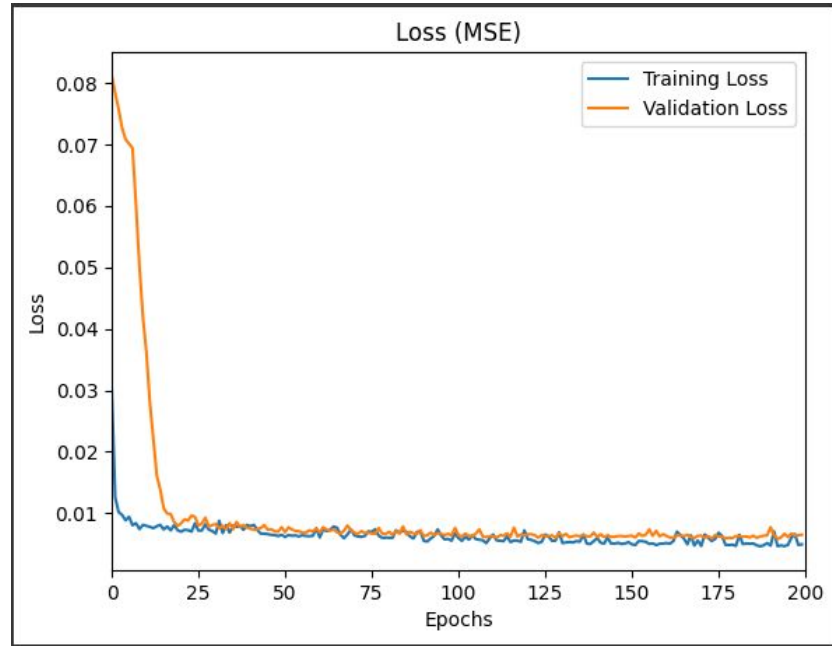
Used Methods/Models

- Learning rate 0.0001



Used Methods/Models

- The U-Net model has a much more stable course of the loss and because skip connection are used the images are sharper




Results

- Baseline model
 - MSE: 0.01020
- U-Net model
 - MSE: 0.00652



Results

- Some predicted images of the U-Net model



Results

- Predicted images of the U-Net model with the lowest MSE score



Results

- Predicted images of the U-Net model with the highest MSE score





Results

- Predicted images on random web images



Results

- Predicted images on random web images with the best MSE



Biggest Problems

- Getting the ImageDataGenerator to work with grayscale images as a input generator and the original label images as a label generator
- Deciding which specific U-Net architecture so use

Outlook

- Possible Improvements
 - Training the model on LAB color representation rather than RGB images to increase performance
 - Using Generative Adversarial Network (GAN)



Thank you